

ECN based Congestion Control for a Software Defined Network

Mohana Prasad S, Manoj Kumar Sure, Krishna M. Sivalingam

Dept. of CSE, Indian Institute of Technology Madras, Chennai, India, 600036

Emails: mohanaprasad1994@gmail.com, skrishnam@iitm.ac.in

Abstract—This paper deals with congestion control in a software defined network (SDN) setting. Presently, explicit router schemes, such as Explicit Congestion Notification (ECN), work in conjunction with the TCP protocol to handle congestion in a distributed manner. With the emergence of SDN and centralized control, it is possible to leverage the global view of the network state to make better congestion control decisions. In this work, we explore the advantages of bringing in global information into distributed congestion control. We propose a framework where the controller with its global view of the network actively participates in the congestion control decisions of the end TCP hosts, by setting the ECN bits of IP packets appropriately. Our framework can be deployed very easily without any change to the end node TCPs or the SDN switches. We also show 30x improvement over the TCP Cubic variant and 1.7x improvement over TCP/RED in terms of flow completion times for one implementation of this framework, using the Mininet emulator.

I. INTRODUCTION

This paper deals with congestion control in computer networks. Existing solutions can be categorized as end-node based or router-based. The latter solutions use queue management and scheduling algorithms that provide signals to the end hosts, to reduce the source traffic. Active Queue Management techniques drop/mark packets at the switch/router buffers thereby signalling the end nodes about congestion [8], [9]. There are some schemes that use both active queue management and end node TCP modifications [6], [10].

Most of the existing congestion control algorithms [3]–[5] have a very limited view of the network and its traffic. Many TCP based congestion control algorithms use packet loss as an indicator of congestion. Another measure used is the round-trip time (RTT). The problem with using RTT is that, the feedback may be easily misinterpreted. For, example consider a 100 packet backlog (with 1,500 Byte packets) in a router queue. It corresponds to 1,200 μ s of queuing delay at 1 Gbps, but only 120 μ s at 10 Gbps. The end node cannot make

fine distinctions without more information. Without the detailed knowledge of the underlying network, TCP will continuously keep increasing and decreasing its congestion window trying to adapt to the network, but may never end up doing so, due to its parochial view.

TCP has been designed to work in a broad range of networks. Each TCP variant works well for some kinds of the network and its traffic and the same TCP performs poorly for other conditions. The interesting part is that, we do not know exactly what objective does TCP congestion control try to optimize [3]. This inflexibility in adapting to new scenarios limits its use.

The emergence of Software Defined Networking (SDN) gives network protocol designers the power of centralized view and centralized control that can be exploited for many applications [13], [14]. SDN provides a centralized view of the network with access to the statistics and other information of the routers and link states. The central controller can aggregate these information and actively participate in the congestion control decisions of the end nodes. The scope of this paper is to study to what extent we can exploit the central view and the centralized control features to improve congestion control in networks.

An SDN-enabled scheme for handling congestion control is presented in the paper. With a global view, the controller knows exactly what each of the link states are and would never misinterpret a packet error as congestion (as was done by many TCP variants). The information at the controller can supplement the indicators like packet loss and delay, that were used by the end nodes earlier. The controller can provide a more realistic view of the network to the end nodes. With a more detailed knowledge about the network and the traffic flowing at any point in time, we can take better, faster congestion control decisions. In the proposed mechanism, the controller instructs the switches (via the OpenFlow API) to set the relevant ECN bits on packets going through a switch. This information is then used by the TCP end-

nodes for changing the TCP congestion window. The scheme has been implemented in the Mininet emulator [15] and studied for three different network scenarios. The results show that the proposed approach achieves improvements over TCP CUBIC and TCP/RED based distributed solutions.

II. BACKGROUND AND RELATED WORK

This section presents the relevant background material and related work.

A. Router-based Congestion Control

The DECbit mechanism [11] is one of the first works that used an explicit congestion control protocol to signal the end nodes about congestion at a router. An router, when it is likely to experience congestion, reacts by marking a bit at the packet header for some of the packets. This bit was then sent by the receiver to the sender for taking appropriate action.

The Random Early Detection (RED) mechanism [9] attempts to maintain an average queue length at the routers. Using threshold values for time-averaged queue lengths, a router drops packets with a probability that increases with the size of the queue. This packet drop results in TCP sources reducing their congestion window and hence transmission rate.

In the Explicit Congestion Notification (ECN) mechanism [2], packets will be marked with Congestion Encountered (CE) bits instead of being dropped. The end-node TCP protocol is modified so that the receiver echoes the CE bits to the sender. When the sender receives such TCP segments, it reacts by reducing the congestion window. TCP's performance can be increased significantly using RED/ECN by setting the parameters with caution. The DCTCP protocol [6] is a variant of TCP optimized for data center networks. DCTCP uses Explicit Congestion Notification (ECN) at the switches and runs its variant of TCP at the end nodes. In [7], another ECN-based scheme that does not require end-host TCP modification is presented for data center networks.

In all these approaches, the decision to drop or mark packets during congestion is taken in a distributed manner by each router. In an SDN-enabled network, it is possible to take advantage of the (logically) centralized control plane to set the bits based on the global view. This paper presents an attempt to explore the advantages of pushing this decision to the SDN controller rather than doing it at the routers.

B. OpenTCP

OpenTCP [1] is a congestion control framework proposed for SDN based networks. The OpenTCP software at the controller collects information about the underlying network such as link utilization values. OpenTCP aims at reducing the Flow Completion Times (FCTs) by updating initial congestion window and retransmission timeout interval for TCP flows. OpenTCP sends Congestion Update Epistles (CUE) packets to the end nodes consisting of information about the suggested changes to their TCP state. OpenTCP's kernel module running in the TCP stack of the end node updates their TCP variant's state using the information in the CUE packets. Thus, OpenTCP requires the end node TCP protocol to be changed.

The approach proposed in this paper is SDN-based and using the ECN mechanism, but does not require any changes to the end nodes' TCP protocol. This is advantageous since changing end nodes may not be possible if all the nodes are not under the same administrator, especially in multi-tenant networks and data centers providing cloud services.

III. PROPOSED SDN-BASED FRAMEWORK

The proposed SDN-enabled framework is shown in Fig. 1. The network consists of several switches/routers that implement the *data plane* with the *control plane* implemented by a logically centralized SDN controller. The end-nodes' TCP protocol implementation is not modified, but is expected to be ECN-enabled. An SDN controller application periodically collects information about the underlying network and the traffic characteristics. Based on these information, it detects congestion according to an algorithm (generically called **Algorithm-A**).

With the global view of the network available, another algorithm (called **Algorithm-B**) selects the end nodes that need to react to any detected congestion. This algorithm also considers the policies and the priorities to flows set by the administrator to decide which end nodes to penalize. After taking these decisions, the application sends new flow rules to the switches (which timeout after a transient amount of time), instructing them to set the ECN bits for particular flows. The application conveys information regarding congestion to the TCP end hosts by setting ECN bits. The normal TCP end nodes, react to packets marked with ECN bits by reducing their congestion window. Different implementations of *Algorithm-A* and *Algorithm-B* can be used for different types of networks and workloads.

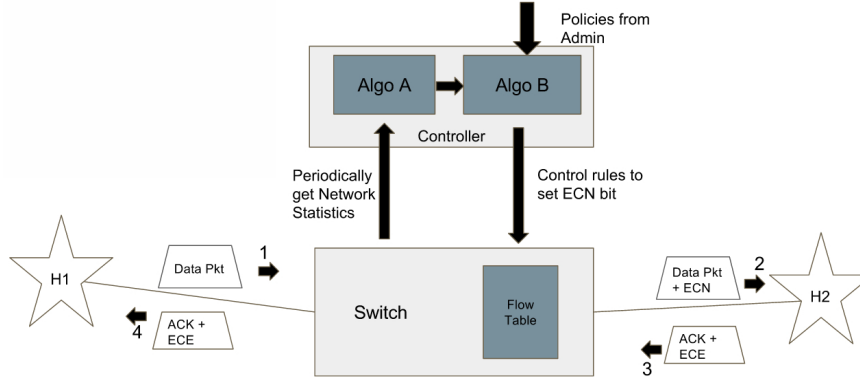


Fig. 1. Proposed SDN-based Congestion Control Framework.

A. Congestion Detection

The congestion detection component (*Algorithm-A*) periodically collects information about the network and its traffic and tries to recognize/predict congestion. SDN switches have the capability of collecting statistics. The application can probe the switch for these statistics and get a complete picture about the network and its present traffic conditions. Queue statistics and link utilization statistics will be used by *Algorithm-A* for detecting congestion. A simple metric to identify potential congestion is if the link utilization or average queue length is greater than a threshold. More sophisticated algorithms could consider global network state and predict congestion based on learning algorithms or other complex heuristics. In the evaluation section, we will present an example for *Algorithm-A*.

B. Handling Congestion

Once *Algorithm-A* predicts congestion in a link/switch, *Algorithm-B* decides on how to handle the congestion. The administrator can also set policies for congestion control. For instance, the policy can specify that certain type of high-priority flows should not be affected. It will compute different ways to avoid congestion, based on the global view and administrator policies. It will identify the set of end nodes whose congestion window can be reduced to avoid congestion while maximizing a global objective.

In traditional TCP, all the end nodes react to congestion by reducing their congestion window. In data center networks and other SDN applications, it is not necessary that all the flows should react to congestion. For instance, we would like the interactive traffic or short-term (i.e. *mice*) flows to be not modified but control the congestion windows of bulk (i.e. *elephant*) traffic. More over, when

all the nodes react to congestion, they all back off at the same time (leading to under utilization) and then move forward simultaneously giving rise to a saw-tooth like bursty traffic at the switches.

For example, the following scheme can be used as *Algorithm-B* in data center networks for prioritizing interactive flows. Whenever congestion is detected in a link, the framework will obtain all statistics about the flows traversing that link. The algorithm can select the top $T\%$ of the flows based on their bandwidth utilization and reduce their congestion windows. Generally, only about 10% of the flows are elephant flows in the network, but they utilize about 90% of the bandwidth. Penalizing the high utilization flows can ensure that interactive traffic is less affected by congestion.

C. Congestion Notification

The next step is to convey the congestion handling information to the end nodes so that they can react appropriately. The *Algorithm-B*, after computing the set nodes that have to reduce their congestion window, will send explicit flow rules to the switches with an additional action of marking the ECN bits apart from forwarding the packets to the right port. These new high priority rules sent by *Algorithm-B* are called the Congestion Control Flow Rules (CC flow rules). When new packets from these flows arrive, they get matched with these top priority CC flow rules and their ECN bits get set at the switch before forwarding.

The ECN bit processing is based on the end-nodes' TCP implementation. As mentioned earlier, we do not require any changes in the TCP implementation. A receiver, upon receiving packets with marked ECN bits, echoes the information to the sender by marking the ECE bit in their ACK packet. The sender acts upon

these marked ACKs and reduce their congestion window according to their native TCP variant.

The new ECN marking flows added are given higher priority in the OpenFlow tables at the switches, compared to the regular flows. This is to make sure they are not skipped. Also, these new CC flow entries are set with a very small rule timeout interval so that adequate number of packets are marked. If a new CC flow rule is added for a very long time, it might send signals of large congestion to the end node that would affect the throughput of that flow adversely. At the same time, the timeout interval should not be too low such that not enough packets from that flow actually get marked. Choosing this timeout value is a critical parameter. Once the timeout interval is over, these extra high priority flows get evicted and the flows start matching with the normal low priority flows rules that were already present in the switches.

Another critical parameter is the periodicity of the network probes. As mentioned earlier, *Algorithm-A* probes the network for network and traffic state. There should be enough time for the system to settle down after the congestion window changes are implemented by the source TCP and the network traffic stabilizes. The ECN packets have to reach the receiver and the receiver has to reply back with bits marked in the ACK packet. It clearly takes about 2 RTT to get communicated to the end nodes and it might take a couple of more RTTs for the network to become stable. Very large probing interval cannot detect sudden congestion and can be less efficient. Thus, the probing interval is a sensitive parameter.

In summary, the salient features and advantages of this framework include:

- 1) **Global View:** This enables the congestion control scheme to obtain a more complete picture about the network state and hence make better congestion predictions that can achieve more globalized objectives. In conventional end-to-end systems, all the end nodes try to optimize their own local objective function and could end up in a Nash equilibrium solution.
- 2) **Prioritization in Congestion control:** Flow priority can be considered in making congestion control decisions and in setting of flow table rules.
- 3) **Fairness:** Fairness can be ensured even across TCP and UDP. Typically, UDP flows can hog the bandwidth placing TCP flows at disadvantage. With the proposed approach, we can monitor the bandwidth used by UDP and use special flow rules to restrict UDP's uneven share of the bandwidth.

- 4) **No change to the end-nodes:** The proposed approach does not require changes to the end-nodes' protocol stack.
- 5) **No change to switches:** This approach uses normal flow rules and the functionalities available with existing SDN switches to achieve congestion control.
- 6) **Easily pluggable CC algorithm:** The Algorithms *A* and *B* can be changed according to different network and traffic needs. A data center network may need a different kind of algorithm than an enterprise network. Congestion control algorithm at the controller becomes a easily changeable mechanism.

IV. PERFORMANCE EVALUATION

This section presents the performance evaluation of the proposed congestion control framework and compared to existing schemes.

A. Implementation Details

The proposed framework has been implemented in an SDN emulator package called Mininet (version 2.2.1) [15]. The Floodlight SDN controller [16], Open vSwitch (version 2.3), OpenFlow 1.3 that supports setting of ECN bits through flow rules have been used. The congestion control framework has been implemented as an application in Floodlight. In all our experiments, the switch contains a single flow table. All the virtual end hosts created by Mininet in our experiments run TCP Cubic implementation that comes with Ubuntu 14.04 kernel. The ECN mechanism has been turned on at all the end nodes. Since Mininet is not able to handle high bandwidths accurately, 100Mbps links have been used everywhere unless stated otherwise.

The analysis is done for data center networks with the objective of achieving lower flow completion times (FCT) for interactive traffic when it co-exists with bulk traffic. As a proof of concept for our framework, we present sample congestion control algorithms, (*Algorithm-A* and *Algorithm-B*), and evaluate their performance.

The proposed *Algorithm-A* probes the switches every 2 seconds and collects port statistics information. From these, we compute the average link utilization of every link in that interval. If the link utilization is greater than 75% , we consider that congestion is likely to occur in that link and inform *Algorithm-B* about it. The handling algorithm, *Algorithm-B*, is defined as follows. Congestion control is imposed on the top $T\%$ of flows

TABLE I
THROUGHPUT OF BULK FLOWS FOR TOPOLOGY 1 (Mbps)

Flow	TCP Cubic	ECN	RED	Proposed
S1	32.5	29.6	30.9	32.7
S2	25.8	29.4	28.1	29.7
S3	26.0	30.2	32.2	28.7
Total	84.3	89.2	91.2	91.1

(based on bandwidth utilization) going through the congested link once the link utilization is greater than 75%. We linearly increase T as link utilization increases and when link utilization reaches 100%, we set $T = 50\%$ penalizing top 50% of the flows when utilization hits the maximum. Since we are not penalizing all the flows at the same time, this can prevent saw tooth like behavior of utilization. Additionally, we start penalizing more and more as we get closer to 100% utilization making sure we have enough bandwidth for new incoming short flows.

The proposed approach is compared with Linux's TCP Cubic scheme, and approaches that use in-network elements in congestion control, namely RED [9] and ECN [2]. The implementation of RED and ECN is available in Mininet.

B. Throughput with long flows

The objective of this experiment is to show that the proposed algorithm does not compromise on bulk flows and achieves good throughput on long-lived traffic.

The topology studied, denoted *Topology 1*, consists of one switch connected to four end nodes with 100 Mbps links each. Three nodes are senders (S_1, S_2, S_3) and one is a receiver (R_1). Three of the senders run *iperf* for two minutes to generate traffic to the receiver. The results are shown in Table I. As seen, the proposed approach gives 8% better total throughput than Cubic and comparable performance with the distributed ECN and RED schemes. The proposed method achieve better fairness than TCP Cubic and is on par with ECN and RED.

The next topology, called *Topology 2*, is a dumbbell topology with 2 switches and 3 nodes connected to each switch. Nodes connected to one of these switches are all senders, called S_1, S_2, S_3 . The nodes connected to the other switch are receivers R_1, R_2, R_3 . Three pairs (S_1, R_1), (S_2, R_2), (S_3, R_3) are selected and long-lived flows are established between them for 2 minutes. All the links have 100 Mbps capacity. The throughput results

TABLE II
THROUGHPUT OF BULK FLOWS FOR TOPOLOGY 2 (Mbps)

Flow	TCP Cubic	ECN	RED	Proposed
S1	27.5	31.3	32.9	31.3
S2	33.3	29.9	30.9	31.9
S3	23.3	32.9	32.1	32.5
Total (Mbps)	84.1	94.1	95.9	95.7

are shown in Table II. As seen, the proposed approach achieves 14% better total throughput than TCP Cubic and is similar to that achieved by RED and ECN. The individual through-puts are also closer to the fair-share values.

C. Coexistence of interactive and bulk traffic

In this set of experiments, we show that the proposed algorithm significantly improves the flow completion time of interactive flows in presence of bulk traffic. In the first experiment, we consider the single switch-four node *Topology 1*. Two long-lived flow between from S_1 and S_2 to the receiver R_1 are established. A 2 MB interactive flow is sent from S_3 to R_1 . This experiment is repeated 30 times for each of the congestion control algorithm and the interactive flow's mean flow completion time is computed, as shown below.

TCP Cubic	ECN	RED	Proposed
10.64	0.36	0.41	0.34

As seen, the proposed scheme obtains 30x improvement to Cubic in flow completion time and about 1.2x improvement to RED. ECN was close (but worse) to our scheme in this topology. We also found that our algorithm shows the least variance in flow completion time compared to the other schemes, based on the 30 iterations.

The *Topology 2*, discussed in the previous section, was next considered and flows established as above. The flow computation time results based on 30 iterations is shown below.

TCP Cubic	ECN	RED	Proposed
10.41	0.46	0.37	0.37

As seen, the proposed approach performs 28x better than TCP Cubic, 1.25x better than ECN, but on par with RED. It was observed that the proposed schemes produce had the least variance among the four.

To summarize, the proposed scheme outperforms TCP Cubic by a large factor. It works the best for both the




Fig. 2. Multi-hop Topology

scenarios, while ECN and RED failed to perform better than the proposed scheme in at least one of the cases.

D. Multi-hop scenario

We evaluate the proposed algorithm in multi-hop networks with multiple bottleneck links. The 7-switch topology is shown in Figure 2. In the first experiment, we establish long lived flows from H_1, H_2, H_4, H_5 to H_7, H_8, H_{10}, H_{11} respectively for 2 minutes. The total throughput (in Mbps) is presented below. As seen, the proposed scheme achieves throughput better than TCP Cubic and comparable to ECN, but less than that of RED.

	TCP Cubic	ECN	RED	Proposed
Throughput	76.7	83.9	85.4	83.8

To check flow completion time for mice flows in this setting, We created long-lived flows from H_1, H_2, H_4, H_5 to H_7, H_8, H_{10}, H_{11} respectively. We simultaneously sent 2MB data from H_3 to H_7 and from H_6 to H_8 and measured the flow completion times. This experiment was repeated 30 times. The average flow completion time (in seconds) achieved by each of the algorithms are shown below.

TCP Cubic	ECN	RED	Proposed
19.46	1.01	1.1	0.64

In multi-hop conditions, our scheme clearly outperforms ECN by a factor of 1.5x and RED by 1.7x without compromising much on throughput of bulk flows. Our scheme also outperforms TCP Cubic by 30x.

V. CONCLUSIONS

This paper has presented a simple and easy to deploy framework for congestion control in Software Defined Networks. The framework is extensible in terms of specific congestion control and detection algorithms. It requires no changes to the end nodes or the SDN-enabled OpenFlow switches. The proposed framework has been

implemented in the Mininet emulator, with heuristics for data center networks. The proposed approach shows 30x improvement to TCP Cubic, 1.7x to RED and 1.5x to ECN on multi-hop topologies for flow completion times of interactive traffic without compromising throughput of bulk flows.

REFERENCES

- [1] M. Ghobadi, S. H. Yeganeh, and Y. Ganjali. "Rethinking end-to-end congestion control in software-defined networks." Proceedings of the 11th ACM Workshop on Hot Topics in Networks (HotNets). 2012.
- [2] S. Floyd. "TCP and explicit congestion notification." ACM SIGCOMM Computer Communication Review 24.5 (1994): 8-23.
- [3] K. Winstein and H. Balakrishnan. "TCP ex machina: Computer-generated congestion control." ACM SIGCOMM Computer Communication Review. Vol. 43. No. 4. ACM, 2013.
- [4] M. Dong, Q. Li, D. Zarchy, P. B. Godfrey and M. Schapira. "PCC: Re-architecting Congestion Control for Consistent High Performance." 12th USENIX Symposium on Networked Systems Design and Implementation (NSDI), 2015.
- [5] J. Perry, A. Ousterhout, H. Balakrishnan, D. Shah, and H. Fugal. "Fastpass: A centralized zero-queue datacenter network". In Proceedings of ACM SIGCOMM, 2014.
- [6] M. Alizadeh, A. Greenberg, D. Maltz, J. Padhye, P. Patel, B. Prabhakar, S. Sengupta, and M. Sridharan. "DCTCP: Efficient Packet Transport for the Commoditized Data Center." In Proceedings of ACM SIGCOMM, 2010.
- [7] H. Wu, J. Ju, G. Lu, C. Guo, Y. Xiong and Y. Zhang. "Tuning ECN for data center networks." Proceedings of the ACM international conference on Emerging networking experiments and technologies (CoNEXT), 2012.
- [8] W.-C. Feng, K. G. Shin, D. Kandlur and D. Saha. "The BLUE active queue management algorithms." IEEE/ACM Transactions on Networking (ToN), 10.4 (2002): 513-528.
- [9] S. Floyd and V. Jacobson. "Random early detection gateways for congestion avoidance." IEEE/ACM Transactions on Networking, 1.4 (1993): 397-413.
- [10] D. Katabi, M. Handley, and C. Rohrs. "Congestion control for high bandwidth-delay product networks." ACM SIGCOMM Computer Communication Review 32.4 (2002): 89-102.
- [11] K. K. Ramakrishnan, and R. Jain. "A binary feedback scheme for congestion avoidance in computer networks." ACM Transactions on Computer Systems (TOCS) 8.2 (1990): 158-181.
- [12] K. K. Ramakrishnan, S. Floyd, and D. Black. "The Addition of Explicit Congestion Notification (ECN) to IP." RFC3168, Sept. 2001
- [13] N. McKeown, T. Anderson, H. Balakrishnan, G. Parulkar, L. Peterson, J. Rexford, S. Shenker, and J. Turner. "OpenFlow: enabling innovation in campus networks." ACM SIGCOMM Computer Communication Review 38, no. 2 (2008): 69-74.
- [14] "Software Defined Networks," <https://www.opennetworking.org/sdn-resources/sdn-definition>, Oct. 2015.
- [15] "Mininet," www.mininet.org, Oct. 2015.
- [16] "Floodlight OpenFlow Controller," www.projectfloodlight.org/floodlight/, Oct. 2015.